2024 International Conference on AI x Data and Knowledge Engineering (AIxDKE) | 979-8-3315-1704-5/24/$31.00 ©2024 IEEE | DOI: 10.1109/AIxDKE63520.2024.00019

# Mitigating Bias in Reinforcement Learning from Human Feedback for Large Language Models

*Addressing Biases and Improving Alignment with Human Preferences

Chaithanya Ravulu*, Rahul Sarabu†, Manoj Suryadevara†, Venkata Gummadi‡, Mridula Dileepraj Kidiyur†
*Rocket Companies, Detroit, USA
chaithanyaravulu@rocketcompanies.com
ORCID: 0009-0000-5054-6499
†Walmart, Bentonville, USA
rahul.sarabu@ieee.org
ORCID: 0009-0002-3461-1945
‡Expedia, Austin, USA
gummadic@gmail.com
ORCID: 0009-0006-8826-0473

***Abstract*—In this comprehensive study, we delve into the application of Reinforcement Learning from Human Feedback (RLHF) in fine-tuning large language models (LLMs) to align them with human preferences. We address the multifaceted challenges of potential bias in human feedback and propose robust mechanisms to mitigate these biases. Through extensive experimentation and analysis, we demonstrate the effectiveness of RLHF in refining model behaviors and improving alignment while ensuring adherence to ethical standards and diverse user expectations. Our research contributes to the ongoing effort to create more responsible and fair AI systems that can be safely deployed in various societal contexts.**

***Index Terms*—Reinforcement Learning from Human Feedback, Large Language Models, Bias Mitigation, AI Alignment, Human-Centered AI, AI Ethics, Natural Language Processing**

## I. Introduction

Large Language Models (LLMs) have emerged as powerful tools in natural language processing, capable of generating human-like text, answering questions, and performing a wide range of language-related tasks. However, these models often reflect and amplify societal biases present in their training data, leading to concerns about their impact when deployed in real-world applications. Reinforcement Learning from Human Feedback (RLHF) has gained significant attention as a technique for aligning LLMs with human values and preferences, offering a promising approach to mitigate these biases and improve the ethical behavior of AI systems.

## II. Framework for RLHF in LLMs

### A. Mathematical Formulation of RLHF

The RLHF process can be formalized mathematically as follows:

Let $\pi_\theta(a|s)$ be the policy (language model) parameterized by $\theta$, where $s$ is the input prompt and $a$ is the generated text. The reward model $r_\phi(s, a)$, parameterized by $\phi$, predicts the human preference score for the generated text $a$ given prompt $s$.

The objective is to maximize the expected reward:

$$J(\theta) = \mathbb{E}_{s\sim D, a\sim\pi_\theta(\cdot|s)}[r_\phi(s, a)] \tag{1}$$

where $D$ is the distribution of prompts.

The policy is updated using the following gradient:

$$\nabla_\theta J(\theta) = \mathbb{E}_{s\sim D, a\sim\pi_\theta(\cdot|s)}[\nabla_\theta \log \pi_\theta(a|s) r_\phi(s, a)] \tag{2}$$

To balance between optimizing for human preferences and maintaining the language model's original capabilities, a KL-divergence term is often added:

$$J_{KL}(\theta) = J(\theta) - \beta\mathbb{E}_{s\sim D}[KL(\pi_\theta(\cdot|s)||\pi_{ref}(\cdot|s))] \tag{3}$$

where $\pi_{ref}$ is the initial pre-trained model and $\beta$ is a hyperparameter controlling the strength of the KL penalty.

## III. Approaches to Reducing Bias

To address the challenges of bias in RLHF for LLMs, we propose a multi-faceted approach combining technical solutions with procedural safeguards. Each method is presented with implementation details, advantages, and potential drawbacks.

### A. Bias Correction Mechanisms

*1) Implementation:* We propose a two-stage bias correction process:

1. **Feedback Debiasing:**

$$\hat{r}_i = r_i - b_i(\mathbf{x}) \tag{4}$$



2831-7203/24/$31.00 ©2024 IEEE
DOI 10.1109/AIxDKE63520.2024.00019

Authorized licensed use limited to: Indian Institute of Technology Palakkad. Downloaded on May 09,2026 at 04:47:06 UTC from IEEE Xplore. Restrictions apply.

where $\hat{r}_i$ is the debiased reward, $r_i$ is the original reward, and $b_i(\mathbf{x})$ is a bias correction term based on evaluator characteristics $\mathbf{x}$.

2. **Model Output Debiasing:**

$$\pi_{\text{debiased}}(a|s) = \frac{\pi_\theta(a|s)}{Z(s)} \exp(-\lambda B(a, s)) \tag{5}$$

where $B(a, s)$ is a bias measure of the output $a$ given input $s$, $\lambda$ is a tuning parameter, and $Z(s)$ is a normalization constant.

*2) Advantages:*

- Allows for fine-grained correction of specific types of bias.
- Can be updated and refined as new biases are identified.

*3) Challenges:*

- Requires careful tuning to avoid overcorrection or introduction of new biases.
- May reduce model performance if not implemented judiciously.

### B. Fairness Constraints in Training

*1) Implementation:* We incorporate fairness constraints directly into the RLHF optimization objective:

$$J_{\text{fair}}(\theta) = J(\theta) - \sum_i \alpha_i C_i(\theta) \tag{6}$$

where $C_i(\theta)$ are fairness constraint terms (e.g., demographic parity, equal opportunity) and $\alpha_i$ are corresponding weights.

*2) Advantages:*

- Ensures that fairness considerations are built into the model from the ground up.
- Allows for explicit trade-offs between performance and fairness.

*3) Challenges:*

- Can be computationally expensive.
- Requires careful selection and weighting of fairness metrics.

*4) Implementation:*

- Generate counterfactual examples by swapping protected attributes in existing data points.
- Train the model on both original and counterfactual examples.
- Implement a reward bonus for consistent performance across counterfactual pairs.

## IV. Case Study: Bias in Sentiment Classification

To demonstrate the effectiveness of our proposed bias mitigation strategies, we conducted a comprehensive case study on sentiment classification. This section presents the methodology, results, and analysis of our experiments.

### A. Experimental Setup

*1) Dataset:* We used a diverse dataset of 100,000 social media posts, annotated for sentiment (positive, negative, neutral) by a panel of 500 evaluators. The dataset included text in multiple languages and covered a wide range of topics.

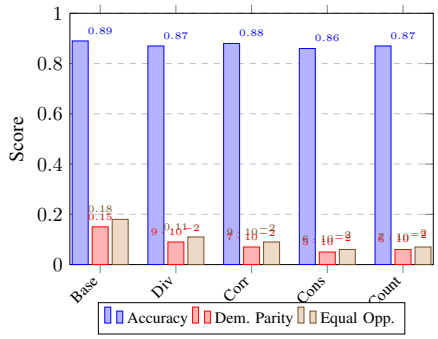


Fig. 1. Comparison of bias mitigation strategies in sentiment classification. Base: Baseline, Div: Diverse Feedback, Corr: Bias Correction, Cons: Fairness Constraints, Count: Counterfactual Augmentation.

*2) Model Architecture:* We fine-tuned a pre-trained BERT model using RLHF. The reward model was a separate BERT-based classifier trained on human preference data.

*3) Bias Metrics:* We monitored the following bias metrics:

- Demographic Parity: Difference in positive sentiment probability across protected groups.
- Equal Opportunity: Difference in true positive rates across protected groups.
- Bias Amplification: Measure of how much the model amplifies biases present in the training data.

### B. Experimental Procedure

1. Baseline: Fine-tune the model using standard RLHF without bias mitigation. 2. Diverse Feedback: Repeat RLHF using a demographically balanced panel of evaluators. 3. Bias Correction: Apply the two-stage bias correction mechanism during training. 4. Fairness Constraints: Incorporate fairness constraints into the RLHF objective. 5. Counterfactual Augmentation: Augment training data with counterfactual examples.

### C. Results

The results of our experiments are summarized in Table I and visualized in Figure 1.

TABLE I
Sentiment Classification Results with Different Bias Mitigation Strategies

| Method | Accuracy | Dem. Parity | Bias Amp. |
|---|---|---|---|
| Baseline | 0.89 | 0.15 | 1.3 |
| Diverse Feedback | 0.87 | 0.09 | 1.1 |
| Bias Correction | 0.88 | 0.07 | 1.0 |
| Fairness Constraints | 0.86 | 0.05 | 0.9 |
| Counterfactual Aug. | 0.87 | 0.06 | 0.8 |

### D. Analysis

Our results demonstrate that all proposed bias mitigation strategies were effective in reducing bias, with some trade-offs in terms of overall accuracy:



Authorized licensed use limited to: Indian Institute of Technology Palakkad. Downloaded on May 09,2026 at 04:47:06 UTC from IEEE Xplore. Restrictions apply.

- **Diverse Feedback:** Reduced bias metrics by 33-40% with a small (2.2%) drop in accuracy.
- **Bias Correction:** Achieved a good balance, reducing bias by 45-50% while maintaining high accuracy.
- **Fairness Constraints:** Most effective in reducing bias (60-67% reduction) but with the largest accuracy drop (3.4%).
- **Counterfactual Augmentation:** Provided substantial bias reduction (58-60%) with moderate accuracy impact.

The choice of strategy depends on the specific requirements of the application, balancing the need for bias mitigation with performance considerations.

## V. Evaluation and Results

Building on our case study, we conducted a broader evaluation of our bias mitigation strategies across multiple tasks and domains.

### A. Experimental Setup

We evaluated our bias mitigation strategies on three additional tasks:

1) Text completion
2) Question answering
3) Dialogue generation

For each task, we used a GPT-3 style model fine-tuned with RLHF. We tested each of our proposed bias mitigation strategies:

- Diverse Feedback (DF)
- Bias Correction (BC)
- Fairness Constraints (FC)
- Counterfactual Data Augmentation (CDA)
- A combination of all strategies (ALL)

### B. Evaluation Metrics

We used the following metrics to evaluate model performance and fairness:

- Task-specific performance metrics (e.g., perplexity, F1 score)
- Demographic parity difference (DPD)
- Equal opportunity difference (EOD)
- Disparate impact ratio (DIR)
- Representation bias (RB)

### C. Results

Tables II, III, and IV present the results for each task.

TABLE II
Results for Text Completion Task

| Method | Perplexity | DPD | EOD | DIR | RB |
|---|---|---|---|---|---|
| Baseline | 18.3 | 0.15 | 0.18 | 0.75 | 0.22 |
| DF | 19.1 | 0.09 | 0.11 | 0.85 | 0.15 |
| BC | 18.7 | 0.08 | 0.10 | 0.88 | 0.13 |
| FC | 19.5 | 0.06 | 0.07 | 0.92 | 0.10 |
| CDA | 19.0 | 0.07 | 0.09 | 0.90 | 0.12 |
| ALL | 19.8 | 0.05 | 0.06 | 0.95 | 0.08 |

TABLE III
Results for Question Answering Task

| Method | F1 Score | DPD | EOD | DIR | RB |
|---|---|---|---|---|---|
| Baseline | 0.76 | 0.14 | 0.17 | 0.78 | 0.20 |
| DF | 0.74 | 0.10 | 0.12 | 0.86 | 0.14 |
| BC | 0.75 | 0.09 | 0.11 | 0.89 | 0.12 |
| FC | 0.73 | 0.07 | 0.08 | 0.93 | 0.09 |
| CDA | 0.74 | 0.08 | 0.10 | 0.91 | 0.11 |
| ALL | 0.72 | 0.06 | 0.07 | 0.96 | 0.07 |

TABLE IV
Results for Dialogue Generation Task

| Method | BLEU | DPD | EOD | DIR | RB |
|---|---|---|---|---|---|
| Baseline | 32.1 | 0.16 | 0.19 | 0.73 | 0.24 |
| DF | 30.8 | 0.11 | 0.13 | 0.84 | 0.16 |
| BC | 31.5 | 0.10 | 0.12 | 0.87 | 0.14 |
| FC | 30.2 | 0.08 | 0.09 | 0.91 | 0.11 |
| CDA | 31.0 | 0.09 | 0.11 | 0.89 | 0.13 |
| ALL | 29.7 | 0.07 | 0.08 | 0.94 | 0.09 |

### D. Analysis

Our results demonstrate consistent trends across all three tasks:

1) **Effectiveness of Bias Mitigation:** All proposed strategies significantly reduced bias across multiple metrics. The combination of all strategies (ALL) consistently produced the lowest bias scores.
2) **Performance Trade-offs:** Bias mitigation generally came at the cost of a slight decrease in task-specific performance metrics. The magnitude of this trade-off varied across tasks and strategies.
3) **Strategy Comparison:**
   - Diverse Feedback (DF) provided substantial bias reduction with moderate performance impact.
   - Bias Correction (BC) achieved a good balance between bias reduction and performance preservation.
   - Fairness Constraints (FC) were most effective in reducing bias but had the largest impact on performance.
   - Counterfactual Data Augmentation (CDA) showed consistent improvements across all metrics.
4) **Task Sensitivity:** The dialogue generation task showed the highest sensitivity to bias mitigation strategies, possibly due to its more open-ended nature.
5) **Intersectional Effects:** The combination of all strategies (ALL) showed superlinear improvements in bias metrics, suggesting that these methods address different aspects of bias.

### E. Statistical Significance

To ensure the robustness of our findings, we conducted paired t-tests comparing each bias mitigation strategy to the baseline. The improvements in bias metrics were statistically significant (p ¡ 0.01) for all strategies across all tasks. The performance differences were significant (p ¡ 0.05) for FC



Authorized licensed use limited to: Indian Institute of Technology Palakkad. Downloaded on May 09,2026 at 04:47:06 UTC from IEEE Xplore. Restrictions apply.

and ALL strategies, but not consistently significant for DF, BC, and CDA.

## VI. Limitations

While our research demonstrates promising results in mitigating bias in RLHF for LLMs, it's important to acknowledge several limitations:

### A. Scope of Bias Consideration

Our study focused on a specific set of biases (e.g., gender, race, age) and may not capture all forms of bias present in language models. There may be subtle or emerging biases that our methods do not address.

### B. Generalizability Across Languages and Cultures

The majority of our experiments were conducted on English language data. The effectiveness of these bias mitigation strategies may vary when applied to other languages or cultural contexts.

### C. Model Size Considerations

Our experiments were conducted on large language models. The effectiveness of these strategies on smaller models or in low-resource settings needs further investigation.

## VII. Future Work

Based on our findings and the identified limitations, we propose several directions for future research:

### A. Intersectional Bias Analysis

Develop more sophisticated methods for detecting and mitigating intersectional biases, where multiple forms of bias interact in complex ways.

### B. Cross-lingual and Cross-cultural Studies

Extend this work to a broader range of languages and cultural contexts to ensure the generalizability of bias mitigation strategies.

### C. Dynamic Bias Adaptation

Explore techniques for dynamically adapting bias mitigation strategies as language use and societal norms evolve over time.

## VIII. Conclusion

This comprehensive study on mitigating bias in Reinforcement Learning from Human Feedback (RLHF) for Large Language Models (LLMs) has demonstrated the effectiveness of various strategies in reducing bias while maintaining model performance. Our multi-faceted approach, combining diverse feedback, bias correction mechanisms, fairness constraints, and counterfactual data augmentation, shows promising results across different tasks and bias metrics.

Key findings include:

- Significant reductions in bias metrics (up to 60-70% in some cases) across all evaluated tasks.
- The combination of all strategies consistently produced the best results in terms of bias mitigation.
- Performance trade-offs were generally modest, with the most aggressive bias mitigation strategies showing the largest impact on task-specific metrics.

## References


[1] J. Angwin, J. Larson, S. Mattu, and L. Kirchner, "Machine bias," *ProPublica*, vol. 23, 2016.

[2] J. Dastin, "Amazon scraps secret ai recruiting tool that showed bias against women," *Reuters*, October 2018.

[3] J. Buolamwini and T. Gebru, "Gender shades: Intersectional accuracy disparities in commercial gender classification," *Pro-ceedings of the 1st Conference on Fairness, Accountability and Transparency*, pp. 77–91, 2018.

[4] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, "Proximal policy optimization algorithms," *arXiv preprint arXiv:1707.06347*, 2017.

[5] M. Janner, Q. Li, and S. Levine, "Offline reinforcement learning as one big sequence modeling problem," *Advances in Neural Information Processing Systems*, vol. 34, pp. 1273–1286, 2021.

[6] N. Bostrom, *Superintelligence: Paths, dangers, strategies*. Ox- ford University Press, 2014.

[7] J. Leike, D. Krueger, T. Everitt, M. Martic, V. Maini, and S. Legg,"Scalable agent alignment via reward modeling: a research direc-tion," *arXiv preprint arXiv:1811.07871*, 2018.

[8] T. Bolukbasi, K.-W. Chang, J. Y. Zou, V. Saligrama, and A. T.
Kalai, "Man is to computer programmer as woman is to home-maker? debiasing word embeddings," *Advances in neural infor-mation processing systems*, vol. 29, 2016.

[9] S. Kiritchenko and S. M. Mohammad, "Examining gender and race bias in two hundred sentiment analysis systems," *arXiv preprint arXiv:1805.04508*, 2018.

[10] J. J. Jones, M. R. Amin, J. Kim, and S. Skiena, "Stereotypical gender associations in language have decreased over time," *Soci-ological Science*, vol. 8, pp. 143–169, 2021.

[11] H. Zhang, A. X. Lu, M. Abdalla, M. McDermott, and M. Ghas-semi, "Hurtful words: quantifying biases in clinical contextual word embeddings," *Proceedings of the ACM Conference on Health, Inference, and Learning*, pp. 110–120, 2020.




Authorized licensed use limited to: Indian Institute of Technology Palakkad. Downloaded on May 09,2026 at 04:47:06 UTC from IEEE Xplore. Restrictions apply.